



NETFLIX RECOMMENDATION SYSTEM

CSYE7200 34322 BIG-DATA SYS ENGR USING SCALA SEC 01 – SPRING 2018

Team - 1

Sonali Chaudhari-001285029
Praneeth Reddy-001225762

OUR PROPOSAL

- **WITH THE HELP OF USER DATA AND RATINGS.**
- **GENERATE A RELEVANT SUGGESTION BASED ON PAST AVAILABLE DATASET.**
- **RECOMMENDING HIGHEST PREDICTED RATING TO A PARTICULAR USER.**
- **MEETING THE DEADLINE OF THE PROJECT.**

USE CASE/ ACTOR

ACTOR

- **USER**
- **THE SYSTEM**

USE CASE

- **USER WILL BELOW OPERATIONS:**
 - **PROVIDING RATING FOR MOVIES**
- **THE APPLICATION WILL PROVIDE LIST OF MOVIES WITH HIGHLY PREDICTED RATINGS**

DETAILS OF DATA

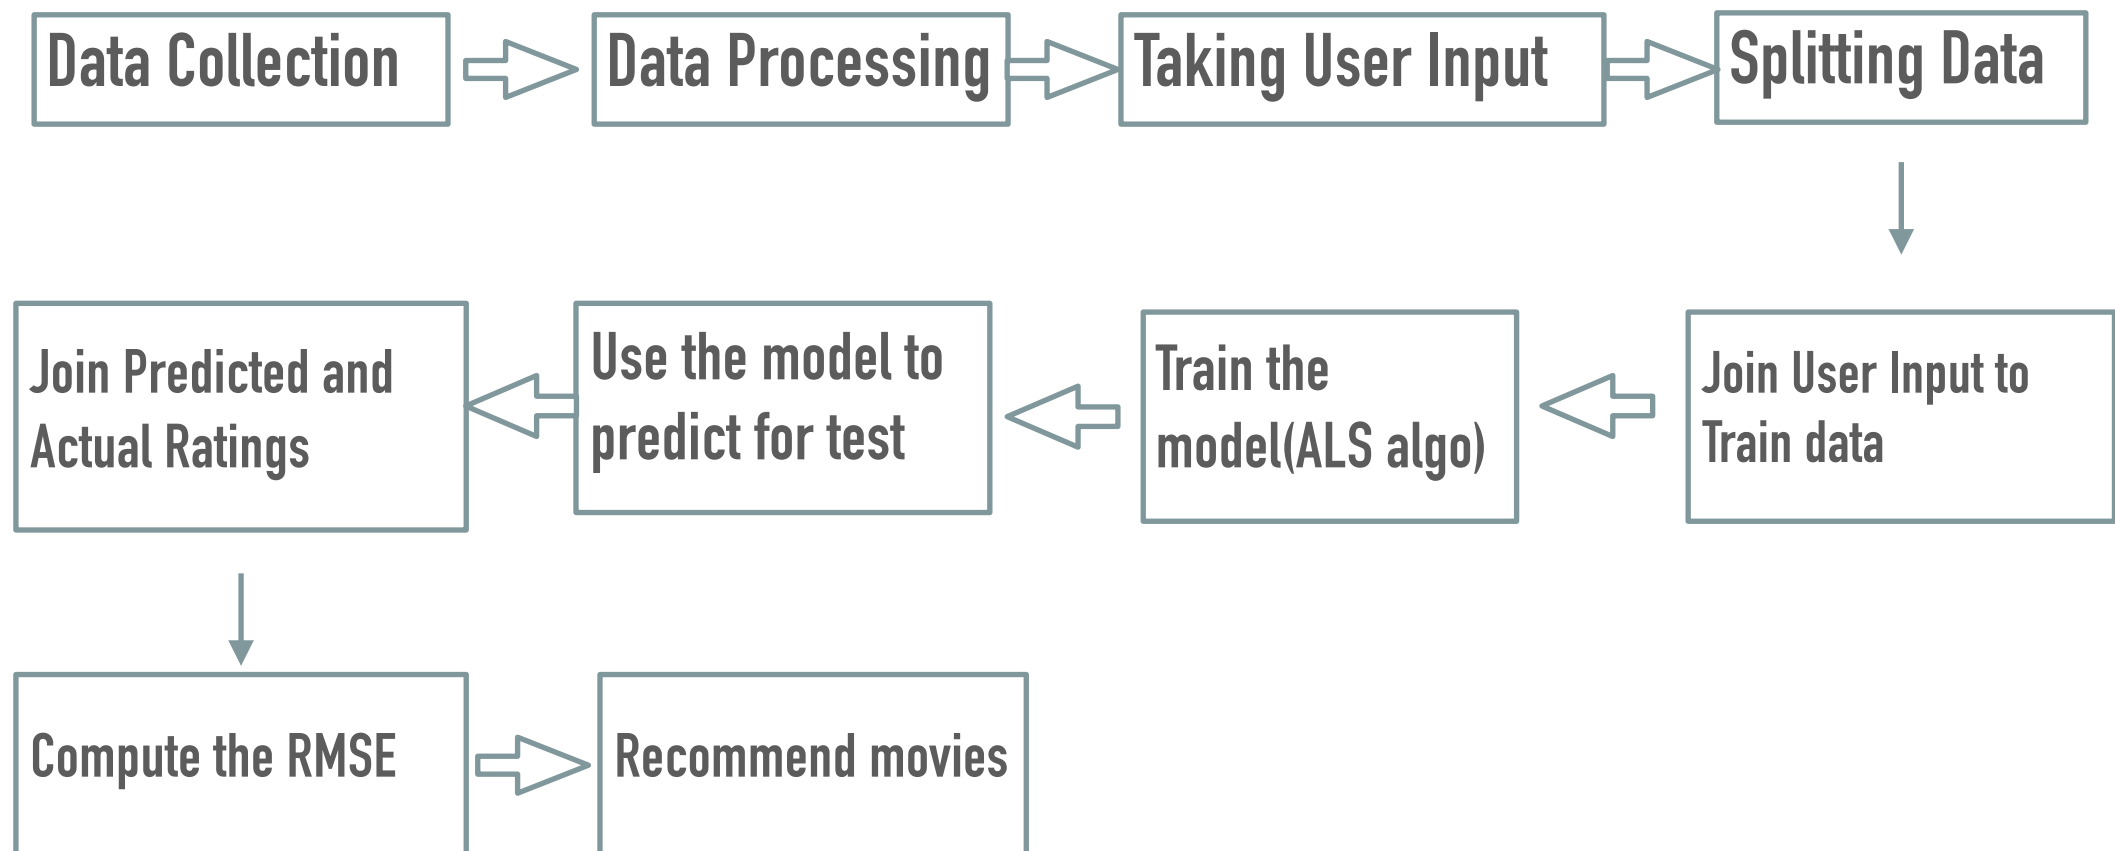
ACTUAL DATA:

- **RATINGS DATA : 10048050**
- **MOVIES : 17770**
- **USER: 480189**

DATA USED (HEAP SIZE ISSUE):

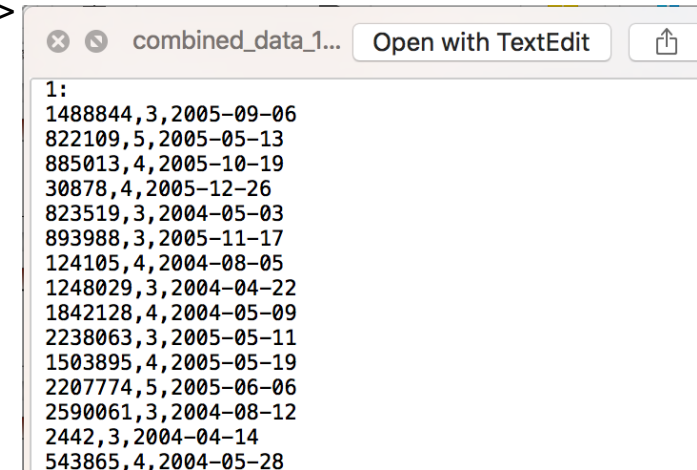
- **RATING DATA : 5010199**
- **MOVIES : 1000**
- **USERS: 404555**

WORKFLOW



DATA CLEANING

- The train data is contained in 4 different text files in the format shown ->

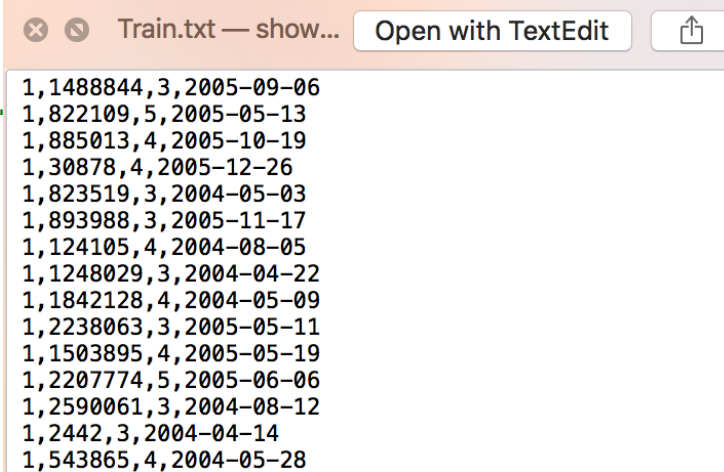


combined_data_1... Open with TextEdit

```
1:  
1488844,3,2005-09-06  
822109,5,2005-05-13  
885013,4,2005-10-19  
30878,4,2005-12-26  
823519,3,2004-05-03  
893988,3,2005-11-17  
124105,4,2004-08-05  
1248029,3,2004-04-22  
1842128,4,2004-05-09  
2238063,3,2005-05-11  
1503895,4,2005-05-19  
2207774,5,2005-06-06  
2590061,3,2004-08-12  
2442,3,2004-04-14  
543865,4,2004-05-28
```

- The four txt file data are converted into the required format and combine into Train.txt file

```
// Generating a combined file in the required format  
val NtflxRecosFile = "/Users/sonalichaudhari/Desktop/netflix-prize-data/"  
val file = new File("/Users/sonalichaudhari/Desktop/netflix-prize-data/Train.txt")  
val bw = new BufferedWriter(new FileWriter(file))  
val train_files = Array("combined_data_1.txt", "combined_data_2.txt", "combined_data_3.txt", "combined_data_4.txt")  
for ( i <- 0 to (train_files.length - 1)) {  
    var app = ""  
    for (line <- Source.fromFile(NtflxRecosFile+train_files(i)).getLines() {  
        if (line.contains(":")) {  
            app = line.toString().stripSuffix(":")  
        }  
        else {  
            var entry = app+","+line  
            bw.write(entry+"\n")  
        }  
    }  
}  
println(("Combined file generated!!"))  
bw.close()
```



Train.txt — show... Open with TextEdit

```
1,1488844,3,2005-09-06  
1,822109,5,2005-05-13  
1,885013,4,2005-10-19  
1,30878,4,2005-12-26  
1,823519,3,2004-05-03  
1,893988,3,2005-11-17  
1,124105,4,2004-08-05  
1,1248029,3,2004-04-22  
1,1842128,4,2004-05-09  
1,2238063,3,2005-05-11  
1,1503895,4,2005-05-19  
1,2207774,5,2005-06-06  
1,2590061,3,2004-08-12  
1,2442,3,2004-04-14  
1,543865,4,2004-05-28
```

TUNING THE MODEL

- ALS algorithm
- Parameters used in the model ALS(rank, iteration, regularization factor)
- Following is how the rise vary with different values of ALS parameters.

ALS(rank, iteration, regularization factor)	RMSE
8,5,0.02	1.10
9, 4, 0.02	1.09
8,10,0.01	1.14
8, 20, 0.01	1.16
10, 5, 0.02	1.09
12, 5, 0.03	1.08
8, 5, 0.05	1.03
8,4,0.05	1.02
8, 4, 0.06	1.015
8, 4, 0.07	1.006
8,5,0.01	1.10
8,5,0.099	0.97

RMSE FOR DIFFERENT RUNS

```
// Model training
val model = ALS.train(training,8,10,0.01)
```

```
// Implementing trained model on the test
```

```
DataProcessing > main(args: Array[String])
```

```
DataProcessing
```

```
18/04/18 18:23:16 INFO spark.SparkContext:
    .scala:111, took 3.795507228 s
```

```
RMSE: 1.141159286583431
```

```
18/04/18 18:23:16 INFO spark
```

```
MatrixFactorizationModel.s
```

```
18/04/18 18:23:16 INFO spark
```

```
shuffle 0 is 162 bytes
```

```
18/04/18 18:23:16 INFO spark
```

```
shuffle 22 is 176 bytes
```

```
18/04/18 18:23:16 INFO spark
```

```
// Model training
```

```
val model = ALS.train(training,8,5,0.01)
```

```
// Implementing trained model on the test
```

```
val prediction = model.predict(test.map(x => (x.user, x.item)))
```

```
// Joining predicted values and actual values
```

```
val predRatings = prediction.map(x => ((x.user, x.item), x.rating))
    .join[Double](test.map(x => ((x.user, x.item), x.rating)))
```

```
DataProcessing > main(args: Array[String])
```

```
DataProcessing
```

```
0.925 s
```

```
18/04/18 18:19:57 INFO spark.SparkContext:
```

```
Job finished: reduce at DataProcessing
```

```
.scala:111, took 3.414699773 s
```

```
RMSE: 1.1098127804418634
```

```
18/04/18 18:19:57 INFO spark.SparkContext:
```

```
Starting job: lookup at
```

```
// Model training
```

```
val model = ALS.train(training,8,5,0.099)
```

```
// Implementing trained model on the test
```

```
val prediction = model.predict(test.map(x => (x.user, x.item)))
```

```
// Joining predicted values and actual values
```

```
DataProcessing > main(args: Array[String])
```

```
DataProcessing
```

```
18/04/18 18:14:25 INFO executor.Executor:
    (TID 165). 945 bytes result sent to driver
```

```
18/04/18 18:14:25 INFO scheduler.TaskSetManager:
    stage 32.0 (TID 165) in 990 ms on localhost
```

```
18/04/18 18:14:25 INFO scheduler.TaskScheduler:
    whose tasks have all completed, from pool
```

```
18/04/18 18:14:25 INFO scheduler.DAGScheduler:
    DataProcessing.scala:111) finished in 0.0
```

```
18/04/18 18:14:25 INFO spark.SparkContext:
    DataProcessing.scala:111, took 3.6811863
```

```
RMSE: 0.9750691117807623
```


PREDICTION ACCURACY

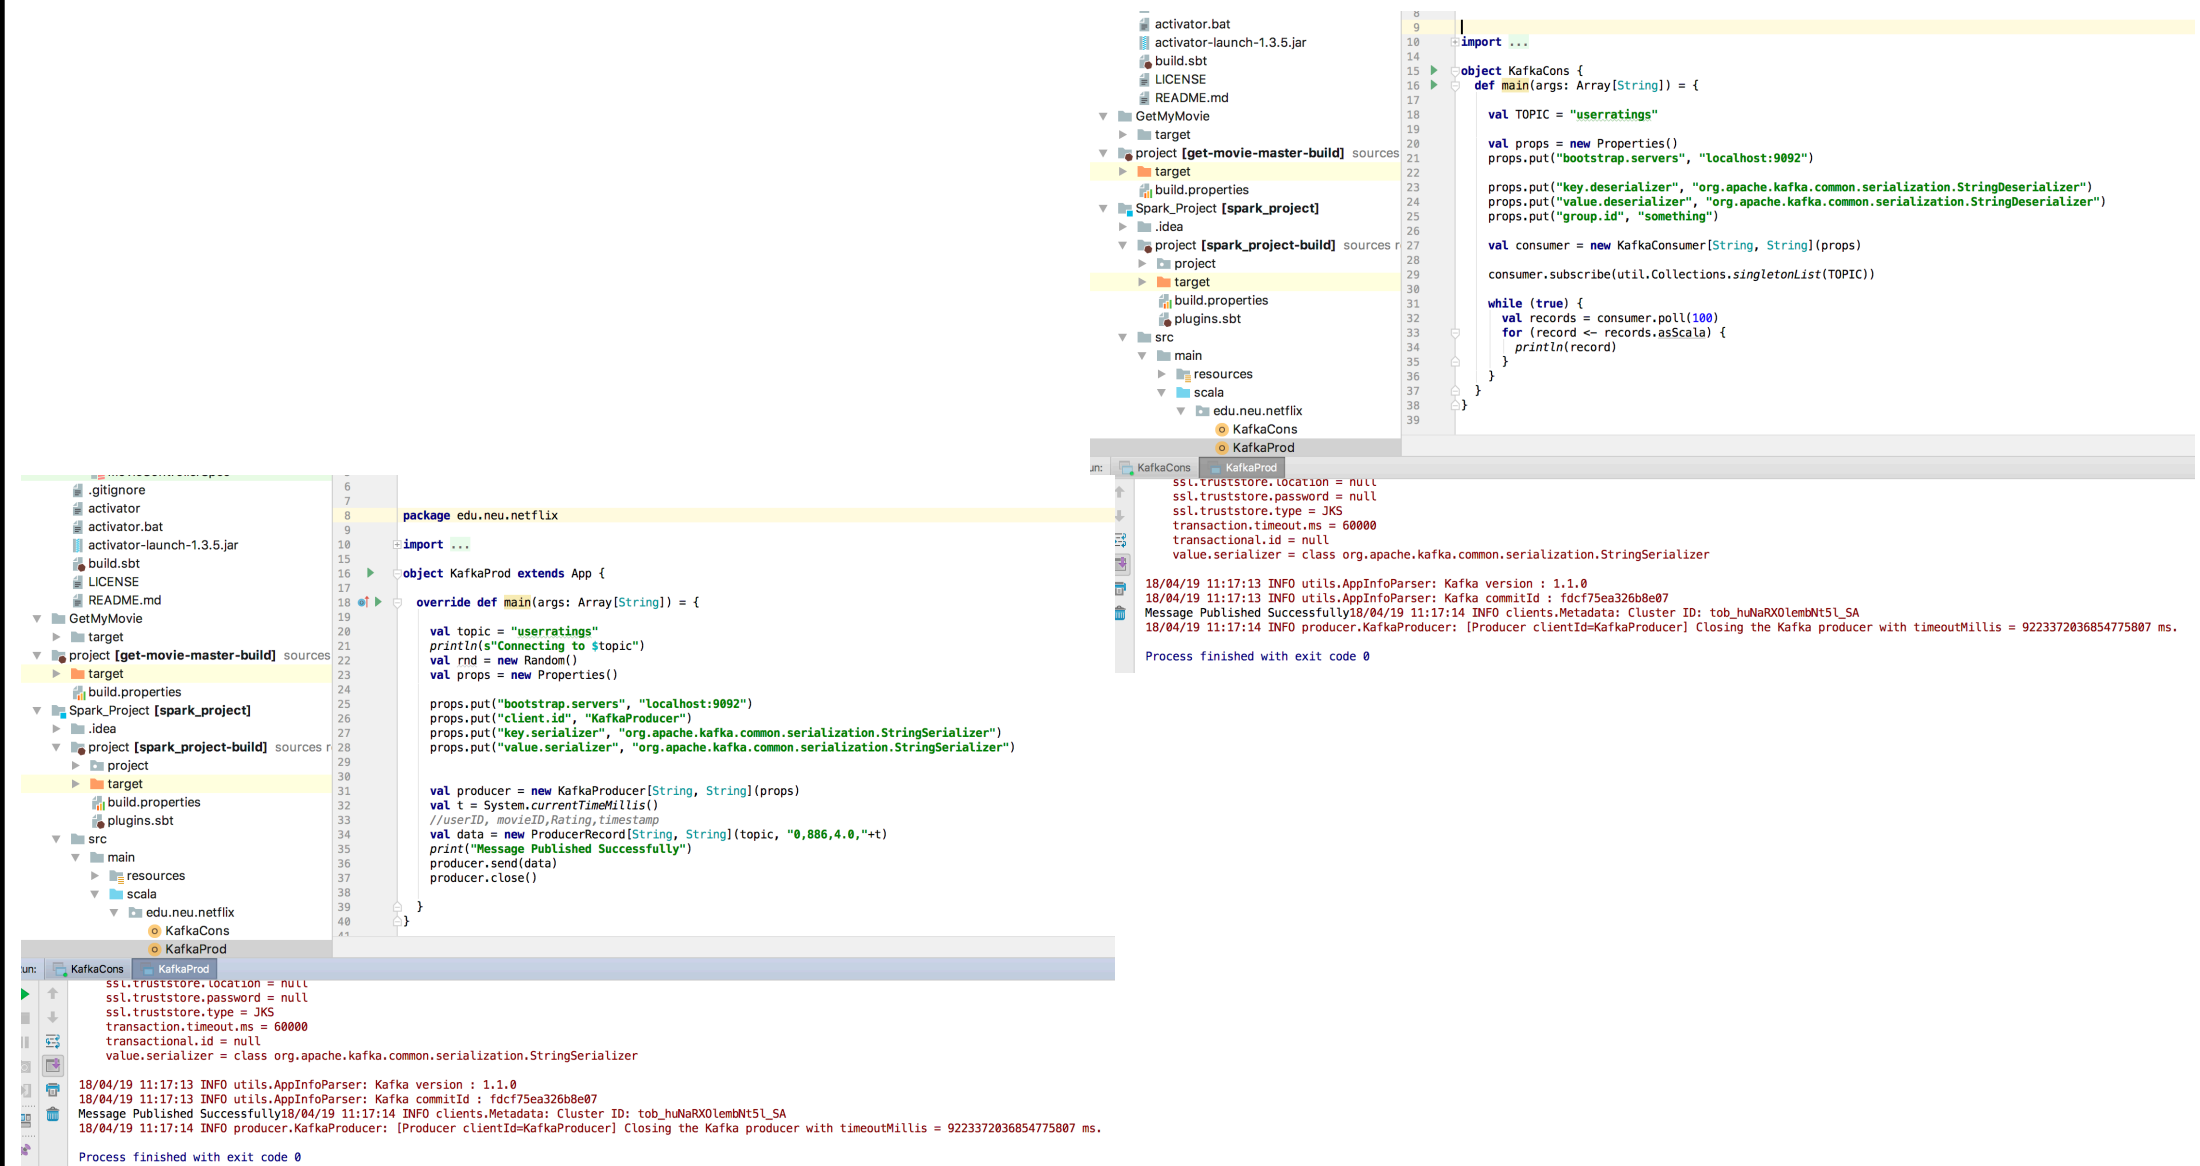
ROOT MEAN SQUARE ERROR (RMSE)

RMSE IS THE PARAMETER USED TO MEASURE THE DIFFERENCE BETWEEN THE PREDICTED VALUES AND THE ACTUAL VALUES.

$$\text{RMSE} = \sqrt{\sum \frac{(y_{pred} - y_{ref})^2}{N}}$$

APPLICATION AND USER INTERFACE

KAFKA PRODUCER AND CONSUMER



The screenshot displays an IDE with two Scala files: `KafkaProd` and `KafkaCons`. The `KafkaProd` file implements a Kafka producer, and the `KafkaCons` file implements a Kafka consumer. Both files are part of the `edu.neu.netflix` package.

```
package edu.neu.netflix

import java.util.Properties
import org.apache.kafka.clients.producer.{KafkaProducer, ProducerRecord}
import org.apache.kafka.common.serialization.StringSerializer

object KafkaProd extends App {
  override def main(args: Array[String]) = {
    val topic = "userratings"
    println(s"Connecting to $topic")
    val rnd = new Random()
    val props = new Properties()

    props.put("bootstrap.servers", "localhost:9092")
    props.put("client.id", "KafkaProducer")
    props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer")
    props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer")

    val producer = new KafkaProducer[String, String](props)
    val t = System.currentTimeMillis()
    //userID, movieID, Rating, timestamp
    val data = new ProducerRecord[String, String](topic, "0.886,4.0," + t)
    print("Message Published Successfully")
    producer.send(data)
    producer.close()
  }
}
```

```
import java.util.Properties
import org.apache.kafka.clients.consumer.{KafkaConsumer}
import org.apache.kafka.common.serialization.StringDeserializer

object KafkaCons {
  def main(args: Array[String]) = {
    val TOPIC = "userratings"

    val props = new Properties()
    props.put("bootstrap.servers", "localhost:9092")

    props.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer")
    props.put("value.deserializer", "org.apache.kafka.common.serialization.StringDeserializer")
    props.put("group.id", "something")

    val consumer = new KafkaConsumer[String, String](props)
    consumer.subscribe(util.Collections.singletonList(TOPIC))

    while (true) {
      val records = consumer.poll(100)
      for (record <- records.asScala) {
        println(record)
      }
    }
  }
}
```

The console output shows the following logs:

```
18/04/19 11:17:13 INFO utils.AppInfoParser: Kafka version : 1.1.0
18/04/19 11:17:13 INFO utils.AppInfoParser: Kafka commitId : fdcf75ea326b8e07
Message Published Successfully18/04/19 11:17:14 INFO clients.Metadata: Cluster ID: tob_huNaRX0lembNt5L_SA
18/04/19 11:17:14 INFO producer.KafkaProducer: [Producer clientId=KafkaProducer] Closing the Kafka producer with timeoutMillis = 9223372036854775807 ms.
Process finished with exit code 0
```

Movie Database

Movie	Rating
Ray	
Speed	
Reservoir Dogs	
Mean Girls	
Something's Gotta Give	
X-Men	
American Beaty	
Rush Hour	
Pay it forward	

<-----Please rate the movies (1(Low) to 5(High)) and Get Suggestions-->

NETFLIX

Show 10 entries

Search:

name	expectedRating
Loading...	

Showing 0 to 0 of 0 entries

Previous

Next

ACCEPTANCE CRITERIA

- **APPLICATION SHOULD BE ABLE TO HANDLE AT LEAST 2500 REQUESTS SIMULTANEOUSLY AND THE MODEL SHOULD BE SCALABLE TO ADD NEW DATA SOURCES AS AND WHEN REQUIRED.**
- **ACHIEVE >90% ACCURACY**

```
// Model training
val model = ALS.train(training,8,5,0.099)

// Implementing trained model on the test
val prediction = model.predict(test.map(x

// Joining predicted values and actual va

DataProcessing > main(args: Array[String])
taProcessing
18/04/18 18:14:25 INFO executor.Executor:
(TID 165). 945 bytes result sent to driv
18/04/18 18:14:25 INFO scheduler.TaskSetM:
stage 32.0 (TID 165) in 990 ms on local
18/04/18 18:14:25 INFO scheduler.TaskSche:
whose tasks have all completed, from poc
18/04/18 18:14:25 INFO scheduler.DAGSched:
DataProcessing.scala:111) finished in 0.
18/04/18 18:14:25 INFO spark.SparkContext
DataProcessing.scala:111, took 3.6811863
RMSE: 0.9750691117807623
```

CHALLENGES FACED

- SPARK, SCALA, KAFKA AND PLAY FRAMEWORK COMPATIBILITY ISSUE
- RESTRUCTURING DATA TO USE INTO CORRECT FORMAT
- OVERFITTING OF THE MODEL
- TUNING THE MODEL
- HEAP SIZE ISSUE

USING PLAY FRAMEWORK

- IMPLEMENTED MVC
- INTEGRATING SPARK
- IMPLEMENTED MOCKITO FOR APPLICATION ,MOVIE
CONTROLLER SPEC TO TEST VARIOUS FEATURES

USING PLAY FRAMEWORK

- **TASK COMPLETED**
 - DATA PROCESSING
 - TAKING USER INPUT
 - PREDICTION GENERATING
- **TASK TO BE COMPLETED**
 - USER INTERFACE

GITHUB REPOSITORY LINK

<https://github.com/reddyse/Big-Data-Engineering-Using-Scala>

THANK YOU...